

JCHARMING: A Directed Model Checking Engine Using Crash Traces to Reproduce Bug.

Mathieu Nayrolles^{*,†,1}, Abdelwahab Hamou-Lhadj¹, Sofiène Tahar² and Alf Larsson³

¹ SBA Research Lab, ECE, Concordia University, Montreal, Canada

² ECE Department, Concordia University, Montreal, Canada

³ PLF System Management, Ericsson, R & D, Stockholm, Sweden

SUMMARY

Due to their inherent complexity, software systems are pledged to be released with bugs. These bugs manifest themselves on client's computers, causing crashes and undesired behaviors. Field crashes, in particular, are challenging to understand and fix as the information provided by the impacted customers are often scarce and inaccurate. To address this issue, there is a need to find ways for automatically reproducing the crash in a lab environment in order to fully understand its root causes. Crash reproduction is also an important step towards developing adequate patches. In this paper, we propose a novel crash reproduction approach, called JCHARMING (Java CrasH Automatic Reproduction by directed Model checkING). JCHARMING uses crash traces and model checking to identify program statements needed to reproduce a crash. Our approach takes advantage of the completeness provided by model checking while ignoring unneeded system states by means of information found in crash traces combined with static slices. We show the effectiveness of JCHARMING by applying it to seven different open source programs cumulating more than one million lines of code scattered in around 7000 classes. Overall, JCHARMING was able to reproduce 85% of the submitted bugs. Copyright © 2015 John Wiley & Sons, Ltd.

Received ...

KEY WORDS: Automatic Bug Reproduction, Dynamic Analysis, Model Checking, Software Maintenance.

1. INTRODUCTION

Despite the efforts spent on testing and verification, there is no guarantee that the final releases of a software system are bug-free, resulting in crashes (field failures) and unwanted behavior of the system. Fixing bugs that cause field crashes is known to be an expensive task. Automated techniques that can help identify the cause of crashes in order to reproduce them can save costs and efforts.

Field failures are, however, challenging to reproduce because the data provided by the end users is often scarce. A survey conducted with developers of major open source software systems such as Apache, Mozilla and Eclipse revealed that one of the most valuable piece of information that can help locate and fix the cause of a crash is the one that can help reproduce it [1]. It is therefore important to invest in techniques and tools for automatic bug reproduction to ease the maintenance process and accelerate the rate of bug fixes and patches.

Existing techniques can be divided into two categories: (a) On-field record and in-house replay [2–4], and (b) In-house crash explanation [5,6]. The first category relies on instrumenting the system

*Correspondence to: Mathieu Nayrolles, SBA Research Lab, ECE, Concordia University, Sir George Williams Campus 1455 De Maisonneuve Blvd. W. Montreal, Quebec, Canada H3G 1M8

†Email: m.nayrol@encs.concordia.ca

in order to capture objects and other system components at run-time. When a faulty behavior occurs in the field, the stored objects as well as the entire heap are sent to the developers along with the faulty methods to reproduce the crash. These techniques tend to be simple to implement and yield good results, but they suffer from two main limitations. First, code instrumentation comes with a non-negligible overhead on the system. The second limitation is that the collected objects may contain sensitive information causing customer privacy issues. The second category is composed of tools leveraging proprietary data in order to provide hints on potential causes. While these techniques are efficient in improving our comprehension of the bugs, they are not designed with the purpose of reproducing them.

In previous work [7] we proposed an hybrid approach called JCHARMING (Java CrasH Automatic Reproduction by directed Model checkING) that uses a combination of crash traces and model checking to automatically reproduce bugs that caused field failures. Unlike existing techniques, JCHARMING does not require instrumentation of the code. It does not need access to the content of the heap either. Instead, JCHARMING uses a list of functions output when an uncaught exception in Java occurs (i.e., the crash trace) to guide a model checking engine to uncover the statements that caused the crash.

Model checking (also known as property checking) is a formal technique for automatically verifying a set of properties of finite-state systems [8]. More specifically, this technique builds a control flow graph where each node represents one state of the program and the set of properties that need to be verified in each state. For real-world programs, model checking is often computationally impracticable because of the state explosion problem [8]. To address this challenge and apply model checking on large programs, we direct the model checking engine towards the crash using program slicing and the content of the crash trace, and hence, reduce the search space.

In our previous work [7], we assessed the effectiveness of using model checking engine directed by stack traces. Indeed, we were capable of achieving 85% accuracy when applied to reproducing bugs from seven different open source systems.

In this extension to our previous work, we improve the performances of our approach by further reducing the number of states to explore and apply JCHARMING to a total of 10 systems (+3), namely Ant, ArgoUML, dnsjava, jfreechart, Log4j, MCT, pdfbox, Mahout, Hadoop and ActiveMQ and 30 bugs (+10) with an accuracy of 91%.

In summary, our contributions from [7] include:

- A first approach, JCHARMING, for reproducing bug using crash traces and directed model checking.
- A directed model checking engine, based on [9], that reduces the number of states to explore in order to reproduce a bug.
- A validation on 7 systems and 20 bugs with an accuracy of 85%.

This paper extends our previous work [7], with the following contributions: an extensive validation of JCHARMING on larger scale systems, *i.e.*, Mahout, Hadoop and, ActiveMQ with 10 more bugs. We further reduce the space search required to reproduce bugs, therefore improving the performances, and improve the accuracy of JCHARMING from 85% to 91%. Furthermore, we provide a more elaborated and up-to-date related work.

The remainder of this paper is organized as follows: In Section 2, we present related work on crash reproduction. In Section 3, we provide some background information on model checking. JCHARMING is the topic of Section 4. Section 5 is dedicated to the case study, followed by threats to validity. We conclude the paper and sketch future directions in Section 8.

2. RELATED WORKS

In this section, we put the emphasis on how stack traces are used in the crash reproduction activity. The published works that fold in this category can be divided into two distinct parts: (A) on field record and in house replay [2,3,10,11] and (B) on house crash explanation [7,12–15].

These two categories yield heterogeneous results depending on the chosen approach and are mainly differentiated by the need of instrumentation. Indeed, the first category need to oversee – by means of instrumentation – the execution of the targeted system on the field in order to reproduce the crashes in house, while tools and approaches belonging to the second category only use materials produced by the crash such as the crash stack or the core dump at crash time. In the first category, tools record different type of evidence such as invoked methods [2], try-catch exception [16] or objects [4]. In the second category, the tools and approaches aim to understand what happens only using material produced by the crash itself, such as crash stack [15], previous – and controlled – execution [14] or model checking [7].

In one hand, the tools and approaches relying on instrumentation face common limitations such as: (A-1) the needs to instrument the source code in order to introduce logging mechanisms [2–4] and therefore (A-2) slowing down the target system. In addition to these limitations, (A-3) recording system behavior by means of instrumentation in order to eventually reproduce on-field crashes yield severe privacy concerns. On the other hand, tools and approaches that only use material that are the result of a crash – such as core dump or exception stack crash – face different limitations. Indeed, in this second category, approaches have to (B-1) reconstruct the timeline of events that have led to the crash [7,15]. Computing all the paths from the initial state of the software to the crash point is a NP-Complex problem and (B-2) approaches struggle with the combinational explosions of states [15,17].

In order to overcome the complexity of the problems, different authors have used different SMT solver [18] and even model checking techniques [19]. However, these techniques require knowledge that fold far behind classical software engineering and are therefore, (B-3) complex to implement [9].

Finally, both categories share a common limitation. Modern softwares are made to interact with the external world and it is possible for the required condition to reproduce a crash to be purely external (AB-4). For example, the reading of a file only present on the hard drive of the customer or the reception of a faulty network packet [7,15].

2.1. On field record and in house replay

Jaygarl *et al.* created OCAT (Object Capture based Automated Testing) [4]. The authors' approach starts by capturing objects created by the program when it runs on-field in order to provide them to an automated test process. The coverage of automated test is often low due to lack of correctly constructed objects. Also, the objects can be mutated by means of evolutionary algorithms. These mutations target primitive field in order to create even more objects and therefore, improve the code coverage. While not directly targeting the reproduction of a bug, OCAT is an approach that was used as the main mechanism for bug reproduction systems.

Narayanasamy *et al.* [2] proposed BugNet, a tool that continuously records program execution for deterministic replay debugging. According to the authors, the size of the recorded data needed to reproduce a bug with high accuracy is around 10MB. This recording is then sent to the developers and allows the deterministic replay of a bug. The authors, argued that with nowadays Internet bandwidth, the size of the recording is not an issue during the transmission of the recorded data. While this approach respects each and every characteristic of the category it is belonging in, it does not provide any responses nor lead to the challenges we have seen in the introduction of this section.

Another approach of this category was proposed by Clause *et al.* [17]. The approach records the execution of the program on the client side and compresses the generated data to the minimal required size that ensures that the reproduction is still achievable. This compression is also done on the client side. Moreover, they keep traces of all accessed documents in the operating system and also compress/reduce them to the minimal. Overall, they are able to reproduce on-field bug using a

file weighting, averaging 70Kb. The minimal execution paths triggering the failure are then sent to the developer who can replay the execution on a sandbox simulating the client's environment. This special feature of the approach proposed by Clause *et al.* [Clause] addresses the AB-4 limitation where crashes are caused by external causes. While the authors broaden the scope of reproducible bug and address the AB-4 limitation, they are recording a lot of private data (A-3). Indeed, all accessed files are saved to enable the reproduction and the operation of the operating system.

Timelapse [20] also address the AB-4 limitation. It focuses on web-application and allows developers to browse and visualize the execution traces recorded by Dolos. Dolos captures and reuses user inputs and network responses to deterministically replay a field crash. Also, both Timelapse and Dolos allow the developers to use conventional tools such as breakpoint and classical debugger. Similarly, private data are recorded without obfuscation of any sort.

Another approach proposed by Shay Artzi *et al.* and named ReCrash records the object states of the targeted programs [3]. They use an in-memory stack which contains every argument and object clones of the real execution in order to reproduce a crash via the automatic generation of unit test cases. Unit test cases are used to provide hints to the developers with buggy code. This approach particularly suffers from the A-2 limitation related to slowing down the execution. Indeed, the overhead for full monitoring is considerably high (between 13% and 64% in some cases). They also propose an alternative in which they record only the methods surrounding the crash. For this work, the crash has to occur at least once so they could use the information causing the crash to identifying the methods surrounding it.

Similarly to ReCrash, JRapture [10] is a capture/replay tool for observation-based testing. The tool captures execution of Java programs to replay it in-house. To capture the execution of a Java program, they used their own version of the Java Virtual Machine (JVM) and employ a lightweight, transparent capture process. Using their own JVM allows capturing any interactions between a Java program and the system, including GUI, files, and console inputs, and on the replay it presents each thread with exactly the same input sequence it saw during capture. However, using a custom JVM is virtually impossible for popular software. Indeed, instead of instrumenting the source code, they require the user to install their JVM that might have some discrepancies with the original one and yield bugs if used with other software. Consequently, JRapture fails at assessing the A-1 and A-3 limitation because it imposes the installation of another JVM that can monitor other software than the intended one.

RECORE (REconstructing CORE dumps) is a tool proposed by Robler *et al.* The tool instruments Java bytecode to wrap every method in a try and catch block while keeping a quasi-null overhead [16]. They start from the core dump and try (with evolutionary algorithms) to reproduce the same dump by executing the programs many times. When the generated dump matches the collected one, they have found the set of input responsible for the failure.

As of now, all presented approaches operate at code level, however, other approaches [11, 21] focus on recording user-GUI interactions. Roehm *et al.* extract the recorder data using delta debugging [22], sequential pattern mining and their combination to reproduce between 75% and 90% of the submitted bug while pruning 93% of the actions.

Among approaches presented here, only the ones proposed by Clause *et al.* and Burg *et al.* address one limitation (AB-4) at the cost, however, of A-3 (privacy). Beside the AB-4 limitation, the RECORE approach has lowered the A-2 limitation (speed). One way of lowering the A-2 and A-3 limitation will be to let user choose where to put the bar between the speed of the targeted program, privacy and bug reproduction efficiency. As an example, users can choose to contribute or not to improving the software – policy employed by many major players such as Microsoft in Visual Studio or Mozilla in Firefox – and propose different type of monitoring where the cost in terms of speed, privacy leaks and efficiency at reproducing bug is clearly explained.

2.2. On house crash explanation

On the other side of the picture, we have the tools and approaches belonging to the on house crash explanation are fewer and newer. Win *et al.* proposed BugRedux for reproducing field failures for in-house debugging [12]. The tool aims to synthesize in-house executions that mimic field failures.

To do so, the authors use several types of data collected in the field such as stack traces, crash stack and point of failure and call sequences. The data that successfully reproduced the field crash is sent to software developers to fix the bug. BugRedux does not particularly assess one limitation and even raise a new one. Indeed, BugRedux lies on several in house executions that are synthesized in order to narrow down the area of research find the crash location and finally, reproduce the bug. However, these in house execution have to be conducted prior to the work on the bug really begins. Also, the in-house executions suffers from the same limitation as unit testing *i.e.*, the executions are based on the developer knowledge and abilities to develop unlikely scenarios and not only almost normal execution. Based on the success of BugRedux, the authors built F3 (Fault localization for Field Failures) [13] and MIMIC [14]. F3 performs many executions of a program on top of BugRedux in order to cover different paths leading to the fault. It then generates many “pass” and “fail” paths which can lead to a better understanding of the bug. They also use grouping, profiling and filtering, to improve the fault localization process. MIMIC further extends F3 by comparing a model of correct behavior to failing executions and identifying violations of the model as potential explanations for failures.

While being close to our approach, BugRedux and F3 may require the call sequence and/or the complete execution trace in order to achieve bug reproduction. When using only the crash traces (referred to as call stack at crash time in their paper), the success rate of BugRedux significantly drops to 37.5% (6/16). The call sequence and the complete execution trace required to reach 100% of bug reproduction can only be obtained through instrumentation and with an overhead ranging from 10% to 1066%. Chronicle [23], is an approach that supports remote debugging by capturing inputs in the application through code instrumentation. They denote themselves by a having a low overhead on the instrumented application, around 10%.

Likewise, Zamfir *et al.* proposed ESD [24], an execution synthesis approach which automatically synthesizes failure execution using only the stack trace information. However, this stack trace is extracted from the core dump and may not always contain the components that caused the crash.

To the best of our knowledge, the most complete work in this category is the one of Chen in his Ph.D thesis [15]. Chen proposed an approach named STAR (Stack Trace used based Automatic crash Reproduction). Using only the crash stack, STAR starts from the crash line and goes backward towards the entry point of the program. During the backward process, STAR computes the required condition using an SMT solver named Yices [18]. The objects that satisfy the required conditions are generated and orchestrated inside a JUnit test case. The test is run and the resulting crash stack is compared to the original one. If both match, the bug is said to be reproduced. Star clearly focuses on assessing the B-2 limitation (state explosion). Indeed, they try to reconstruct the events in a backward fashion and therefore save numerous states to explore. Also, Star do well with B-3 (complexity to implement) as it uses Yices [18] and potentially Z3 [25] (stated in future work) that are well-known and well-supported SMT solver.

Except for STAR, approaches targeting the reproduction of field crashes require the instrumentation of the code or the running platform in order to save the stack call or the objects to successfully reproduce bugs. As we discussed earlier, instrumentation can cause a massive overhead (1% to 1066%) while running the system. In addition, data generated at run-time using instrumentation may contain sensitive information.

JCHARMING’s directed model checking overcome the state explosion problem (B-2) of classical model checking and allow the generation of Junit test case in a reasonable amount of time. In order to address the limitation B-1 (reconstruct the events), B-2 (state explosion) and B-3 (complex implementation) authors have adopted different approaches such as execution synthesis, backward analysis, well-known and supported SMT and directed model checking, respectively. While all approaches assess B-1 and are able to yield acceptable results, it seems that no good trade-off can be found between the simplicity of implementation (B-3) and the reduction of state explosion (B-2).

3. PRELIMINARIES

Model checking (also known as property checking) will, given a system (that could be software [19] or hardware based [26]), check if the system meets a specification *Spec* by testing exhaustively all the states of the system under test (SUT), which can be represented by a Kripke [27] structure:

$$SUT = \langle S, T, P \rangle \quad (1)$$

where *S* is the set of states, $T \subseteq S * S$ represents the transitions between the states and *P* is the set of properties that each state satisfies. The SUT is said to satisfy a set of properties *p* when there exists a sequence of states transition *x* leading towards these properties. This can be written as:

$$(SUT, x) \models p \quad (2)$$

However, this only ensures that $\exists x$ such that *p* is reached at some point in the execution of the program and not that *p* holds nor that $\forall x, p$ is satisfiable. In JCHARMING, SUTs are bound to a simple specification: they must not crash under a fair environment. In the framework of this study, we consider a fair environment as any environment where the transitions between the states represent the functionalities offered by the program. For example, in a fair environment, the program heap or other memory spaces cannot be modified. Without this fairness constraint, all programs could be tagged as buggy since we could, for example, destroy objects in memory while the program continues its execution. As we are interested in verifying the absence of unhandled exceptions in the SUT, we aim to verify that for all possible combinations of states and transitions there is no path leading towards a crash. That is:

$$\forall x. (SUT, x) \models \neg c \quad (3)$$

If such a path exists (i.e., $\exists x$ such that $(SUT, x) \models c$) then the model checker engine will output the path *x* (known as the counter-example) which can then be executed. The resulting Java exception crash trace is compared with the original crash trace to assess if the bug is reproduced. While being accurate and exhaustive in finding counter-examples, model checking suffers from the state explosion problem, which hinders its applicability to large software systems.

To show the contrast between testing and model checking, we use the hypothetical example of Figures 1, 2 and 3 to sketch the possible results of each approach. These figures depicts a toy program where from the entry point, unknown calls are made (dotted points) and, at some points, two methods are called. These methods, called `Foo.Bar` and `Bar.Foo`, implement a for loop from 0 to `loopCount`. The only difference between these two methods is that the `Bar.Foo` method throws an exception if *i* becomes larger than two. Hereafter, we denote this property as $c_{i>2}$.

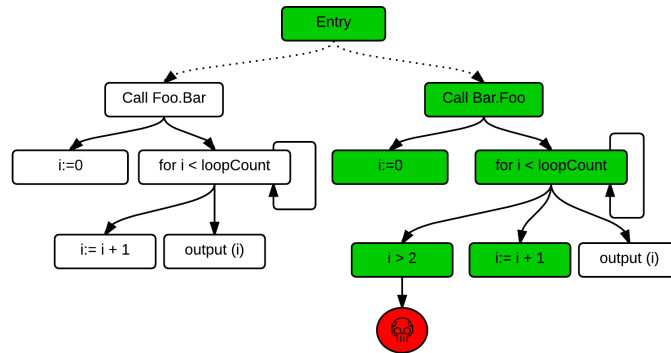


Figure 1. A toy program under testing

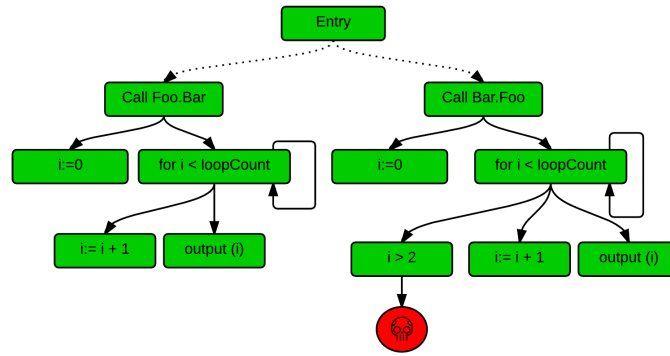


Figure 2. A toy program under model checking

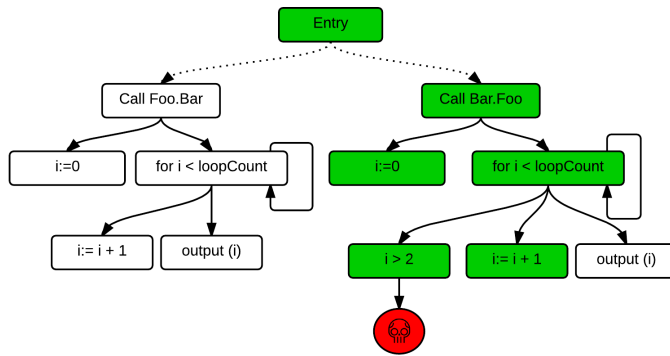


Figure 3. A toy program under directed model checking

Figure 1 shows the program statements that could be covered using testing approaches. Testing software is a demanding task where a set of techniques is used to test the SUT according to some input.

Software testing depends on how well the tester understands the SUT in order to write relevant test cases that are likely to find errors in the program. Program testing is usually insufficient because it is not exhaustive. In our case, using testing would mean that the tester knows what to look for in order to detect the causes of the failure. We do not assume this knowledge in JCHARMING.

Model checking, on the other hand, explores each and every state of the program (Figure 2), which makes it complete, but impractical for real-world and large systems. To overcome the state explosion problem of model checking, directed (or guided) model checking has been introduced [28]. Directed model checking use insightsgenerally heuristicabout the SUT in order to reduce the number of states that need to be examined. Figure 3 explores only the states that may lead to a specific location, in our case, the location of the fault. The challenge, however, is to design techniques that can guide the model checking engine. As we will describe in the next section, we use crash traces and program slicing to overcome this challenge.

4. THE JCHARMING APPROACH

Figure 4 shows an overview of JCHARMING. The first step consists of collecting crash traces, which contain raw lines displayed to the standard output when an uncaught exception in Java occurs. In the second step, the crash traces are preprocessed by removing noise (mainly calls to Java standard library methods). The next step is to apply backward slicing using static analysis to

expand the information contained in the crash trace while reducing the search space. The resulting slice along with the crash trace are given as input to the model checking engine. The model checker executes statements along the paths from the main function to the first line of the crash trace (i.e., the last method executed at crash time, also called the crash location point). Once the model checker finds inconsistencies in the program leading to a crash, we take the crash stack generated by the model checker and compare it to the original crash trace (after preprocessing). The last step is to build a JUnit test, to be used by software engineers to reproduce the bug in a deterministic way.




Figure 4. Overview of JCHARMING.

4.1. Collecting Crash Traces

The first step of JCHARMING is to collect the crash trace caused by an uncaught exception. Crash traces are usually included in crash reports and can therefore be automatically retrieved using a simple regular expression. Figure 5 shows an example of a crash trace that contains the exception thrown when executing the program depicted in Figures 1 to 3. The crash trace contains a call to the `Bar.foo()` method the crash location point and calls to Java standard library functions (in this case, GUI methods because the program was launched using a GUI).

```

1.java.lang.InvalidActivityException:loopTimes
should be < 3
2. at Foo.bar(Foo.java:10)
3. at GUI.buttonActionPerformed(GUI.java:88)
4. at GUI.access0(GUI.java : 85)
5.atGUI1.actionPerformed(GUI.java:57)
6. caused by java.lang.IndexOutOfBoundsException : 3
7. at jsep.Foo.buggy(Foo.java:17)
8. and 4 more ...

```

Figure 5. Java `InvalidActivityException` is thrown in the `Bar.foo()` loop if the control variable is greater than 2.

As shown in Figure 5, we can see that the first line (referred to as frame f_0 , subsequently the next line is called frame f_1 , etc.) does not represent the real crash point but it is only the last exception of a chain of exceptions. Indeed, the `InvalidActivity` has been triggered by an `IndexOutOfBoundsException` in `jsep.Foo.buggy`. This kind of crash traces reflects several nested try/catch blocks.

In addition, it is common in Java to have incomplete crash traces. According to the Java documentation [29], line 8 of Figure 5 should be interpreted as follows: “This line indicates that the remainder of the stack trace for this exception matches the indicated number of frames from the bottom of the stack trace of the exception that was caused by this exception (the “enclosing exception”). This shorthand can greatly reduce the length of the output in the common case where a wrapped exception is thrown from the same method as the “causative exception” is caught.”

We are likely to find shortened traces in bug repositories as they are what the user sees without any possibility to expand their content.

4.2. Preprocessing

In the preprocessing step, we first reconstruct and reorganize the crash trace in order to address the problem of nested exceptions. Then, with the aim to obtain an optimal guidance for our directed model checking engine, we remove frames that are out of our control. Frames out of our controls refer usually, but are not limited to, Java library methods and third party libraries. In Figure 5, we can see that Java GUI and event management components appear in the crash trace. We assume that these methods are not the cause of the crash; otherwise it means that there is something wrong with the on- field JDK. If this is the case, we will not be able to reproduce the crash. Note that removing these unneeded frames will also reduce the search space of the model checker.

4.3. Building the Backward Static Slice

For large systems, a crash trace does not necessary contain all the methods that have been executed starting from the entry point of the program (i.e., the main function) to the crash location point. We need to complete the content of the crash trace by identifying all the statements that have been executed starting from the main function until the last line of the preprocessed crash trace. In Figure 5, this will be the function call `Bar.foo()`, which happens to be also the crash location point. To achieve this, we turn to static analysis by extracting a backward slice from the main function of the program to the `Bar.foo()` method.

A backward slice contains all possible branches that may lead to a point n from a point m as well as the definition of the variables that control these branches [30]. In other words, the slice of a program point n is the program subset that may influence the reachability of point n starting from point m . The backward slice containing the branches and the definition of the variables leading to n from m is noted as $bslice_{[m \leftarrow n]}$.

We perform a static backward slice between each frame to compensate for possible missing information in the crash trace. More formally, the final static backward slice is represented as follows:

$$bslice_{[entry \leftarrow f_0]} = bslice_{[f_1 \leftarrow f_0]} \cup bslice_{[f_2 \leftarrow f_1]} \cup \dots \cup bslice_{[f_n \leftarrow f_{n-1}]} \cup bslice_{[entry \leftarrow f_n]} \quad (4)$$

Note that the union of the slices computed between each pair of frames must be a subset of the final slice between f_0 and the entry point of the program. More formally:

$$\bigcup_{i=0}^{entry} bslice_{[f_{i+1} \leftarrow f_i]} \subseteq bslice_{[entry \leftarrow f_0]} \quad (5)$$

Indeed, in Figure 6, the set of states allowing to reach f_0 from f_2 is greater than the set of states to reach f_1 from f_2 plus set of states to reach f_0 from f_1 . In this hypothetical example and assuming that z_2 is a prerequisite to f_2 then $bslice_{[entry \leftarrow f_0]} = \{f_0, f_1, f_2, z_0, z_1, z_2, z_3\}$ while $\bigcup_{i=0}^n bslice_{[f_{i+1} \leftarrow f_i]}$.

In the worst case scenerio where there exists one and only one transition between each frame, which is very unlikely for real and complex systems, then $bslice_{[entry \leftarrow f_0]}$ and $\bigcup_{i=0}^n bslice_{[f_{i+1} \leftarrow f_i]}$ yield the same set of states with a comparable computational cost since the number of branches to explore will be the same in both cases.

Algorithm 1 is a high level representation of how we compute the backward slice between each frame. The algorithm takes as input the pre-processed call trace, the byte code of the SUT, and the entry point. From line 1 to line 5, we initialize the different variables used by the algorithm. The main loop of the algorithm begins at line 6 and ends at line 15. In this loop, we compute the static slice between the current frame and the next one. If the computed static slice is not empty then we




Figure 6. Hypothetical example representing $bslice_{[entry \leftarrow f_0]}$ Vs. $\cup_{i=0}^n bslice_{[f_{i+1} \leftarrow f_i]} = \{f_0, f_1, f_2, z_2\}$

update the final backward slice with the newly computed slice.

Data: Crash Stack, BCode, Entry Point

Result: BSolve

Frames frames \leftarrow extract frames from crash stack;

Int *n* \leftarrow size of frame;

Int *offset* \leftarrow 1;

Bslice *bSlice* \leftarrow \emptyset ;

for *i* \leftarrow 0 to *i* $<$ *n* && *offset* $<$ *n* - 1 **do**

 BSlice *currentBSlice* \leftarrow backward slice from frames[*i*] to *i* + *offset*;

if *currentBSlice* $\neq \emptyset$ **then**

bSlice \leftarrow *bSlice* \cup *currentBSlice*;

offset \leftarrow 1;

else

offset \leftarrow *offset* + 1;

end

end

Algorithm 1: High level algorithm computing the union of the slices

Using backward slicing, the search space of the model checker that processes the example of Figures 1 to 3 is given by the following expression:

$$\exists x. \left(\bigcup_{i=0}^{entry} bslice_{[f_{i+1} \leftarrow f_i]} \subset SUT \right) \models c_{i>2} \quad (6)$$

That is, there exists a sequence of states transitions x that satisfies $c_{i>2}$ where both the transitions and the states are entry elements of $\bigcup_{i=0}^{entry} bslice_{[f_{i+1} \leftarrow f_i]}$. Obviously, $c_{i>2}$ also needs to be included for the final static slice to be usable by the model checking engine. Consequently, the only frame that need to be untouched for the backward static slice to be meaningful is f_0 .

4.4. Directed Model Checking

The model checking engine we use in this paper is called JPF (Java PathFinder) [9], which is an extensible JVM for Java bytecode verification. This tool was first created as a front-end for the SPIN model checker [31] in 1999 before being open-sourced in 2005. JPF is organized around five simple operations: (i) *generate states*, (ii) *forward*, (iii) *backtrack*, (iv) *restore state* and (v) *check*. In the forward operation, the model checking engine generates the next state s_{t+1} . If s_{t+1} has successors then it is saved in a backtrack table to be restored later. The backtrack operation consists of restoring the last state in the backtrack table. The restore operation allows restoring any state and can be used to restore the entire program as it was the last time we choose between two

branches. After each, forward, backtrack and restore state operation the check properties operation is triggered.

In order to direct JPF, we have to modify the *generate states* and the *forward* steps. The *generate states* is populated with entry the states in $\bigcup_{i=0}^{entry} bslice_{[f_{i+1} \leftarrow f_i]} \subset SUT$ and we adjust the *forward step* to explore a state if the target state $s_i + 1$ and the transition x to pass from the current state s_i to s_{i+1} are in $\bigcup_{i=0}^{entry} bslice_{[f_{i+1} \leftarrow f_i]} \subset SUT$ and $x. \bigcup_{i=0}^{entry} bslice_{[f_{i+1} \leftarrow f_i]} \subset x.SUT$.

4.5. Validation

To validate the result of directed model checking, we modify the *check properties* step that checks if the current sequence of states transitions x satisfies a set a property. If the current states transitions x can throw an exception, we execute x and compare the exception thrown to the original crash trace (after preprocessing). If the two exceptions match, we conclude that the conditions needed to trigger the failure have been met and the bug is reproduced.

However, as argued by Kim et al. in [32], the same failure can be reached from different paths of the program. Although the states executed to reach the defect are not exactly the same, they might be useful to enhance the understanding of the bug by software developers, and speed up the deployment of a fix. Therefore, in this paper, we consider a defect to be partially reproduced if the crash trace generated from the model checker matches the original crash trace by a factor of t , where t is a threshold specified by the user. t is the percentage of identical frames between both crash traces.

4.6. Generating Test Cases for Bug Reproduction

To help software developers reproduce the crash in a lab environment we automatically produce the JUnit test cases necessary to run the SUT to cause the exercise of the bug.

To build a test suite that reproduces a defect, we need to create a set of objects used as arguments for the methods that will enable us to travel from the entry point of the program to the defect location. JPF has the ability to keep track of what happens during model checking in the form of traces containing the visited states and the value of the variables. We leverage this capability to create the required objects and call the methods leading to the failure location.

During the testing of the SUT, JPF emits a trace that is composed of the executed instructions. For industrial size system, this trace can contain mullions of instructions. This trace is stored in memory and therefore can be queried while JPF is running. However, accessing the trace during the JPF execution considerably slows down the checking process as both the querying mechanism and the JPF engine compete with each other for resources. In order to allow JPF to use 100% of the available resources and still be able to query the executed instruction post-mortem, we implemented a `listener` that listens the JPF trace emitter. Each time on new instruction is processed by JPF, our listener catch it and then save it to a mongodb database to be queried post-mortem. Figure 7 presents a high level architecture of the different components.

When the `validate` step trigger a crash stack similarity superior to a factor t , the Junit generation engine queries the mongoDB database and fetches the sequence of instructions that led to the crash of interest. Figure 7 contains an hypothetical sequence of instructions related to the example of Figures 1, 2, 3 which reads : `new jsme.Bar, invokespecial jsme.Bar(), astore_1 [bar], aload_1 [bar], iconst_3, invokevirtual jsme.Bar.foo(int), const_0, istore_2 [i], goto, iload_2 [i], iconst_2, if_icmple iinc 2 1, new java.lang.Exception`. From this sequence we know that, to reproduce to crash of interest, we have to (1) create new object `jsme.Bar` (`new jsme.Bar, invokespecial jsme.Bar()`), (2) store the newly created object in a variable named `bar` (`astore_1 [bar]`), (3) invoke the method `jsme.Bar.foo(int)` of the `bar` object with 3 as value (`aload_1 [bar], iconst_3, invokevirtual jsme.Bar.foo(int)`). Then, the `jsme.Bar.foo(int)` method will executes a loop from 0 to 3 and throw an exception at `i = 3` (`const_0, istore_2 [i], goto, iload_2 [i], iconst_2, if_icmple iinc 2 1, new java.lang.Exception`).

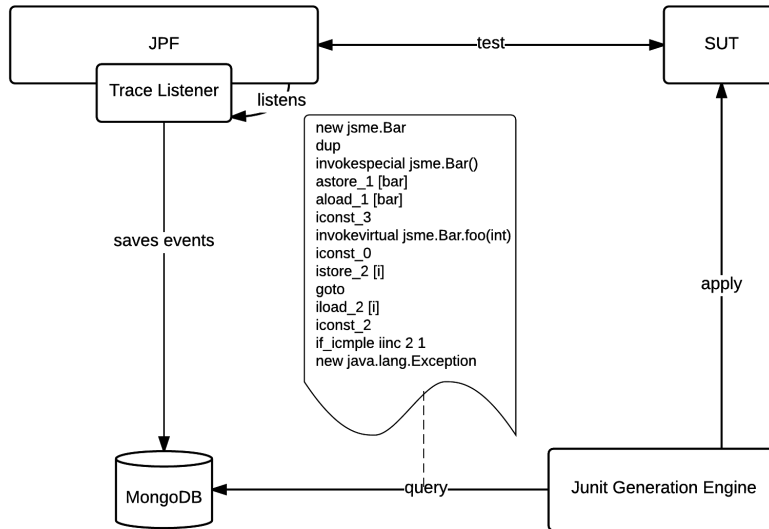


Figure 7. High level architecture of the JUnit test case generation

The generation of the JUnit test itself is based on templates and targets directly the system under test. Templates are excerpts of Java source code with well-defined tags that will be replaced by values. We use templates because the JUnit test case have common structures as shown by Figure 8.

```

public class {% SUT %}-{% BUG %} extends TestCase {

    private static final String faillure = {% CRASH_STACK %};
    private static final int threshold = {% THRESHOLD %};
    private static int differences = Integer.MAX_VALUE;
    private static final StringTokenizer tokenizerFaillure =
        new StringTokenizer(faillure , "\n");

    @Test
    public test{% SUT %}() {
        try {
            {% STEPS %}
        } catch (Exception e) {
            // Count the differences
        }
        assertTrue(differences <= tokenizeOriginalFaillure
            .countTokens() / 100 * (threshold - 100));
    }
}
  
```

Figure 8. Simplified Unit Test template

5. CASE STUDIES

In this section, we show the effectiveness of JCHARMING to reproduce bugs in seven open source systems[†]. The aim of the case study is to answer the following question: *Can we use crash traces and directed model checking to reproduce on- field bugs in a reasonable amount of time?*

5.1. Targeted Systems

Table I shows the systems and their characteristics in terms of Kilo Line of Code (KLoC) and Number of Classes (NoC).

Table I. List of target systems in terms of Kilo line of code (KLoC), number of classes (NoC) and Bug # ID

SUT	KLOC	NoC	Bug #ID
Ant	265	1,233	38622, 41422
ArgoUML	58	1,922	2603, 2558, 311, 1786
dnsjava	33	182	38
jfreechart	310	990	434, 664, 916
Log4j	70	363	11570, 40212, 41186, 45335, 46271, 47912, 47957
MCT	203	1267	440ed48
pdfbox	201	957	1,412, 1,359
Mahout	183	1,242	
Hadoop	62	6,337	
ActiveMQ	132	3,797	
Total	1,517	17,348	30

Apache Ant [35] is a popular command-line tool to build make files. While it is mainly known for Java applications, Apache Ant also allows building C and C++ applications. We choose to analyze Apache Ant because it has been used by other researchers in similar studies.

ArgoUML [36] is one of the major players in the open source UML modeling tools. It has many years of bug management and, similar to Apache Ant, it has been extensively used as a test subject in many studies.

Dnsjava [37] is a tool for the implementation of the DNS mechanisms in Java. This tool can be used for queries, zone transfers, and dynamic updates. It is not as large as the other two, but it still makes an interesting case subject because it has been well maintained for the past decade. Also, this tool is used in many other popular tools such as Aspirin, Muffin and Scarab.

JfreeChart [38] is a well-known library that enables the creation of professional charts. Similar to dnsjava, it has been maintained over a very long period of time JfreeChart was created in 2005 and it is a relatively large application.

Apache Log4j [39] is a logging library for Java. This is not a very large library, but it is extensively used by thousands of programs. As other Apache projects, this tool is well maintained by a strong open source community and allows developers to submit bugs. The bugs which are in the bug report system of Log4j are, generally speaking, well documented and almost every bug contains a related crash trace and, therefore, it is a tool of interest to us.

MCT [40] stands for Mission Control technologies and was developed by the NASA Ames Research Center (the creators of JPF) for use in spaceflight mission operation. This tool benefits from two years of history and targets a very critical domain, Spacial Mission Control. Therefore, this tool has to be particularly and carefully tested and, consequently, the remaining bugs should be hard to discover and reproduce.

[†]The bug reports used in this study and the result of the model checker are made available for download from research.mathieu-nayrolles.com/jcharming/

PDFBox [41] is another tool supported by the Apache Software Foundation since 2009 and was created in 2008. PDFBox allows the creation of new PDF documents and the manipulation of existing documents.

5.2. Bug Selection and Crash Traces

In this study, we have selected the reproduced bugs randomly in order to avoid the introduction of any bias. We selected a random number of bugs ranging from 1 to 10 for each SUT containing the word “exception” and where the description of the bug contains a match a regular expression designed to find the pattern of a Java exception.

6. RESULTS

Table II shows the results of JCHARMING in terms of Bug #ID, reproduction status, and execution time (in minutes) of directed model checking (DMC) and Model Checking (MC). The experiments have been conducted on a Linux machine (8 GB of RAM and using Java 1.7.0_51).

- The result is noted as “Yes” if the bug has been fully reproduced, meaning that the crash trace generated by the model checker is identical to the crash trace collected during the failure of the system.
- The result is “Partial” if the similarity between the crash trace generated by the model checker and the original crash trace is above $t=80\%$. Given an 80% similarity threshold, we consider partial reproduction as successful. A different threshold could be used.
- Finally, the result of the approach is reported as “No” if either the similarity is below $t ; 80\%$ or the model checker failed to crash the system given the input we provided.

Table II. Effectiveness of JCHARMING using directed model checking (DMC) and model checking (MC) in minutes

SUT	Bug #ID	Reprod.	Time DMC	Time MC
Ant	38622	Yes	25.4	-
	41422	No	-	-
ArgoUML	2558	Partial	10.6	-
	2603	Partial	9.4	-
	311	Yes	11.3	-
	1786	Partial	9.9	-
DnsJava	38	Yes	4	23
jFreeChart	434	Yes	27.3	-
	664	Partial	31.2	-
	916	Yes	26.4	-
Log4j	11570	Yes	12.1	-
	40212	Yes	15.8	-
	41186	Partial	16.7	-
	45335	No	-	-
	46271	Yes	13.9	-
	47912	Yes	12.3	-
	47957	No	-	-
MCT	440ed48	Yes	18.6	-
PDFBox	1412	Partial	19.7	-
	1359	No	-	-

As we can see in Table II, we were able to reproduce 17 bugs out of 20 bugs either completely or partially (85ratio). The average time to reproduce a bug is 16 minutes. This result demonstrates the

effectiveness of our approach, more particularly, the use of backward slicing to create a manageable search space that guides adequately the model checking engine. We also believe that our approach is usable in practice since it is also time efficient. Among the 20 different bugs we have tested, we will describe one bug (chosen randomly) for each category (successfully reproduced, partially reproduced, and not reproduced) for further analysis.

6.1. Successfully reproduced

The first bug we describe in this discussion is the bug #311 belonging to ArgoUML. This bug was submitted in an earlier version of ArgoUML. This bug is very simple to manually reproduce thanks to the extensive description provided by the reporter, which reads: *“I open my first project (Untitled Model by default). I choose to draw a Class Diagram. I add a class to the diagram. The class name appears in the left browser panel. I can select the class by clicking on its name. I add an instance variable to the class. The attribute name appears in the left browser panel. I can’t select the attribute by clicking on its name. Exception occurred during event dispatching:”*

The reporter also attached the following crash trace that we used as input for JCHARMING:

```

1. java.lang.NullPointerException:
2. at
3. uci.uml.ui.props.PropPanelAttribute .setTargetInternal (PropPanelAttribute.java)
4. at uci.uml.ui.props.PropPanel. setTarget(PropPanel.java)
5. at uci.uml.ui.TabProps.setTarget(TabProps.java)
6. at uci.uml.ui.DetailsPane.setTarget (DetailsPane.java)
7. at uci.uml.ui.ProjectBrowser.select (ProjectBrowser.java)
8. at uci.uml.ui.NavigatorPane.mySingleClick (NavigatorPane.java)
9. at uci.uml.ui.NavigatorPane$Navigator MouseListener.mouse Clicked(NavigatorPane.java)
10.at java.awt.AWTEventMulticaster.mouseClicked (AWTEventMulticaster.java:211)
11. at java.awt.AWTEventMulticaster.mouseClicked (AWTEvent Multicast er.java:210)
12.at java.awt.Component.processMouseEvent (Component.java:3168)
...
19. java.awt.LightweightDispatcher .retargetMouseEvent (Container.java:2068)
22. at java.awt.Container .dispatchEventImp l(Container.java:1046)
23. at java.awt.Window .dispatchEventImpl (Window.java:749)
24. at java.awt.Component .dispatchEvent (Component.java:2312)
25. at java.awt.EventQueue .dispatchEvent (EventQueue.java:301)
28. at java.awt.EventDispatchThread.pumpEvents
(EventDispatch Thread.java:90)
29. at java.awt.EventDispatchThread.run(EventDispatch Thread.java:82)

```

The cause of this bug is that the reference to the attribute of the class was lost after being displayed on the left panel of ArgoUML and therefore, selecting it through a mouse click throws a null pointer exception. In the subsequent version, ArgoUML developers added a TargetManager to keep the reference of such object in the program. Using the crash trace, JCHARMING’s preprocessing step removed the lines between lines 11 and 29 because they belong to the Java standard library and we do not want neither the static slice nor the model checking engine to verify the Java standard library but only the SUT. Then, the third step performs the static analysis following the process described in Section IV.C. The fourth step performs the model checking on the static slice to produce the same crash trace. More specifically, the model checker identifies that the method `setTargetInternal(Object o)` could receive a null object that will result in a Null pointer exception.

6.2. Partially reproduced

As an example of a partially reproduced bug, we explore the bug #664 of the Jfreechart program. The description provided by the reporter is: *“In ChartPanel.mouseMoved there’s a line of code*

which creates a new *ChartMouseEvent* using as first parameter the object returned by *getChart()*. For *getChart()* is legal to return null if the chart is null, but *ChartMouseEvent*'s constructor calls the parent constructor which throws an *IllegalArgumentException* if the object passed in is null."

The reporter provided the crash trace containing 42 lines and the replaced an unknown number of lines by the following statement "ideleted entryi". While JCHARMING successfully reproduced a crash yielding almost the same trace as the original trace, the "ideleted entryi" statement – which was surrounded by calls to the standard java library – was not suppressed and stayed in the crash trace. That is, JCHARMING produced only the 6 (out of 7) first lines and reached 83% similarity, and thus a partial reproduction.

```
1. java.lang.IllegalArgumentException: null source
2. at java.util.EventObject.<init> (EventObject.java:38)
3. at
4. org.jfree.chart.ChartMouseEvent.<init> (ChartMouseEvent.java:83)
5. at org.jfree.chart.ChartPanel .mouseMoved(ChartPanel.java:1692)
6. <deleted entry>
```

In all bugs that were partially reproduced, we found that the differences between the crash trace generated from the model checker and the original crash trace (after preprocessing) consists of few lines only.

6.3. Not Reproduced

To conclude the discussion on the case study, we present a case where JCHARMING was unable to reproduce the failure. For the bug #47957 belonging to Log4j and reported in late 2009 the reporter wrote: "*Configure SyslogAppender with a Layout class that does not exist; it throws a NullPointerException. Following is the exception trace:*" and attached the following crash trace:

```
1. 10052009 01:36:46 ERROR [Default: 1] struts.CPEExceptionHandler.execute
RID[(null;25KbxIK0voima4h00ZLBQFC;236A18E60000045C3A 7D74272C4B4A61)]
2. Wrapping Exception in ModuleException
3. java.lang.NullPointerException
4. at org.apache.log4j.net.SyslogAppender .append(SyslogAppender.java:250)
5. at org.apache.log4j.AppenderSkeleton .doAppend(AppenderSkeleton.java:230)
6. at org.apache.log4j.helper.AppenderAttachableImpl .appendLoopOnAppen-
ders(AppenderAttachableImpl .java:65)
7. at org.apache.log4j.Category.callAppenders (Category.java:203)
8. at org.apache.log4j.Category .forcedLog(Category.java:388)
9. at org.apache.log4j.Category.info (Category.java:663)
```

The first three lines are not produced by the standard execution of the SUT but by an *ExceptionHandler* belonging to Struts [42]. Struts is an open source MVC (Model View Controller) framework for building Java Web Application. JCHARMING examined the source code of Log4J for the crash location *struts.CPEExceptionHandler.execute* and did not find it since this method belongs to the source base of Struts – which uses log4j as a logging mechanism. As a result, the backward slice was not produced, and we failed to perform the next steps. It is noteworthy that the bug is marked as duplicate of the bug #46271 which contains a proper crash trace. We believe that JCHARMING could have successfully reproduced the crash, if it was applied to the original bug.

7. THREATS TO VALIDITY

The selection of SUTs is one of the common threats to validity for approaches aiming to improve the understanding of program's behavior. It is possible that the selected programs share common properties that we are not aware of and therefore, invalidate our results. However, the SUTs analyzed by JCHARMING are the same as the ones used in similar studies. Moreover, the SUTs vary in terms of purpose, size and history.

Another threat to validity lies in the way we have selected the bugs used in this study. We selected the bugs randomly to avoid any bias. One may argue that a better approach would be to select bugs based on complexity or other criteria (severity, etc.). We believe that a complex bug (if complexity can at all be measured) may perhaps have an impact on the running time of the approach, but we are not convinced that the accuracy of our approach depends on the complexity or the type of bugs we use. Instead, it depends on the quality of the produced crash trace.

In addition, we see a threat to validity that stems from the fact that we only used open source systems. The results may not be generalizable to industrial systems. The last author of the paper expressed an interest to apply these techniques to Ericsson systems. We intend to undertake these studies in future work.

Field failures can also occur due to the running environment on which the program is executed. For instance, the failure may have been caused by the reception of a network packet or the opening of a given file located on the hard drive of the users. The resulting failures will hardly be reproducible by JCHARMING.

Finally, the programs we used in this study are all based on the Java programming language and JCHARMING leverages the crash traces produced by the JVM to reproduce bugs. This can limit the generalization of the results. However, similar to Java, .Net, Python and Ruby languages also produce crash traces. Therefore, JCHARMING could be applied to other object-oriented languages.

In conclusion, internal and external validity have both been minimized by choosing a relatively large set of different systems and using input data that can be found in other programming languages.

8. CONCLUSION AND FUTURE WORK

We presented JCHARMING (Java CrasH Automatic Reproduction by directed Model checking), an automatic bug reproduction technique that combines crash traces and directed model checking. JCHARMING takes advantage of the completeness provided by model checking while ignoring unwanted states by using the content of crash traces augmented with a backward slice. When applied to twentybugs from six open source systems, JCHARMING was able to reproduce fully or partially 85% of the bugs.

To build on this work, we need to experiment with additional (and more complex) bugs with the dual aim to (a) improve and fine tune the approach, and (b) assess the scalability of our approach when applied to even larger (and proprietary) systems. Finally, we want to test the performance of JCHARMING on bugs due to multi-threading.

ACKNOWLEDGEMENT

This class file was developed by Sunrise Setting Ltd, Paignton, Devon, UK. Website:
www.sunrise-setting.co.uk

REFERENCES

1. Bettenburg N, Just S, Schröter A, Weiss C, Premraj R, Zimmermann T. What makes a good bug report? *Proceedings of the 16th ACM SIGSOFT International Symposium on Foundations of software engineering - SIGSOFT '08/FSE-16*, ACM Press: New York, New York, USA, 2008; 308, doi:10.1145/1453101.1453146. URL <http://portal.acm.org/citation.cfm?doid=1453101.1453146>.
2. Narayanasamy S, Pokam G, Calder B. BugNet: Continuously Recording Program Execution for Deterministic Replay Debugging. *Proceedings of the 32nd annual international symposium on Computer Architecture*, vol. 33,

- ACM, 2005; 284–295, doi:10.1145/1080695.1069994. URL <http://dl.acm.org/citation.cfm?id=1080695.1069994>.
3. Artzi S, Kim S, Ernst MD. Recrash: Making software failures reproducible by preserving object states. *Proceedings of the 22nd European conference on Object-Oriented Programming*, 2008; 542–565.
 4. Jaygarl H, Kim S, Xie T, Chang CK. OCAT: Object Capture based Automated Testing. *Proceedings of the 19th international symposium on Software testing and analysis*, 2010; 159–170.
 5. Manevich R, Sridharan M, Adams S. PSE: explaining program failures via postmortem static analysis. *ACM SIGSOFT Software Engineering Notes*, vol. 29, ACM, 2004; 63, doi:10.1145/1041685.1029907. URL <http://dl.acm.org/citation.cfm?id=1041685.1029907>.
 6. Chandra S, Fink SJ, Sridharan M. Snugglebug: a powerful approach to weakest preconditions. *ACM Sigplan Notices*, vol. 44, ACM, 2009; 363–374.
 7. Nayrolles M, Hamou-Lhadj A, Sofiene T, Larsson A. JCHARMING : A Bug Reproduction Approach Using Crash Traces and Directed Model Checking. *SANER'15*, 2015; 101–110.
 8. Baier C, Katoen Jp. *Principles of Model Checking*. MIT press Cambridge, 2008.
 9. Visser W, Pseanu CS, Khurshid S. Test input generation with java PathFinder. *Proceedings of the 2004 ACM SIGSOFT international symposium on Software testing and analysis - ISSTA '04* 2004; :97doi:10.1145/1007512.1007526. URL <http://portal.acm.org/citation.cfm?doid=1007512.1007526>.
 10. Steven J, Chandra P, Fleck B, Podgurski A. jRapture: A Capture/Replay Tool for Observation-Based Testing. *Proceedings of the International Symposium on Software Testing and Analysis*, August, 2000; 158–167.
 11. Roehm T, Nosovic S, Bruegge B. Automated extraction of failure reproduction steps from user interaction traces. *2015 IEEE 22nd International Conference on Software Analysis, Evolution, and Reengineering (SANER)*, IEEE, 2015; 121–130, doi:10.1109/SANER.2015.7081822. URL <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=7081822>.
 12. Jin W, Orso A. BugRedux: Reproducing field failures for in-house debugging. *2012 34th International Conference on Software Engineering (ICSE)*, Ieee, 2012; 474–484, doi:10.1109/ICSE.2012.6227168. URL <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=6227168>.
 13. Jin W, Orso A. F3: fault localization for field failures. *Proceedings of the 2013 International Symposium on Software Testing and Analysis - ISSTA 2013*, ACM Press: New York, New York, USA, 2013; 213–223, doi:10.1145/2483760.2483763. URL <http://dl.acm.org/citation.cfm?doid=2483760.2483763>.
 14. Zuddas D, Jin W, Pastore F, Mariani L, Orso A. MIMIC: locating and understanding bugs by analyzing mimicked executions. *Proceedings of the 29th ACM/IEEE international conference on Automated software engineering - ASE '14*, ACM Press: New York, New York, USA, 2014; 815–826, doi:10.1145/2642937.2643014. URL <http://dl.acm.org/citation.cfm?id=2642937.2643014>.
 15. Chen N. Star: stack trace based automatic crash reproduction. PhD Thesis 2013.
 16. Rößler J, Zeller A, Fraser G, Zamfir C, Candea G. Reconstructing Core Dumps. *Proceedings of the 6th International Conference on Software Testing, Verification and Validation, ser. ICST*, 2013.
 17. Clause J, Orso A. A Technique for Enabling and Supporting Debugging of Field Failures. *ICSE '07 Proceedings of the 29th international conference on Software Engineering*, 2007; 261–270.
 18. Dutertre B, Moura LD. The yices smt solver. *Technical Report* 2006.
 19. Visser W, Havelund K, Brat G, Park S, Lerda F. Model Checking Programs. *Automated Software Engineering*, vol. 10, Springer, 2003; 203–232.
 20. Burg B, Bailey R, Ko AJ, Ernst MD. Interactive record/replay for web application debugging. *Proceedings of the 26th annual ACM symposium on User interface software and technology - UIST '13*, ACM Press: New York, New York, USA, 2013; 473–484, doi:10.1145/2501988.2502050. URL <http://dl.acm.org/citation.cfm?id=2501988.2502050>.
 21. Herbold S, Grabowski J, Waack S, Bünting U. Improved Bug Reporting and Reproduction through Non-intrusive GUI Usage Monitoring and Automated Replaying. *2011 IEEE Fourth International Conference on Software Testing, Verification and Validation Workshops*, IEEE, 2011; 232–241, doi:10.1109/ICSTW.2011.66. URL <http://ieeexplore.ieee.org/articleDetails.jsp?arnumber=5954414>.
 22. Zeller A, Hildebrandt R. Simplifying and isolating failure-inducing input. *IEEE Transactions on Software Engineering* 2002; **28**(2):183–200, doi:10.1109/32.988498. URL <http://dl.acm.org/citation.cfm?id=506201.506206>.
 23. Bell J, Sarda N, Kaiser G. Chronicler: Lightweight recording to reproduce field failures. *2013 35th International Conference on Software Engineering (ICSE)*, IEEE, 2013; 362–371, doi:10.1109/ICSE.2013.6606582. URL <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=6606582>.
 24. Zamfir C, Candea G. Execution Synthesis: A Technique for Automated Software Debugging. *Proceeding of EuroSys '10 Proceedings of the 5th European conference on Computer systems*, 2010; 321–334.
 25. De Moura L, Bjørner N. Z3: An efficient SMT solver. *Tools and Algorithms for the Construction and Analysis of Systems*. Springer, 2008; 337–340.
 26. Kropf T. *Introduction to formal hardware verification*. Springer, 1999.
 27. Kripke SA. Semantical Considerations on Modal Logic. *Acta Philosophica Fennica* 1963; **16**(1963):83–94.
 28. Rungta N, Mercer EG. Guided model checking for programs with polymorphism. *Proceedings of the 2009 ACM SIGPLAN workshop on Partial evaluation and program manipulation - PEPM '09*, ACM Press: New York, New York, USA, 2009; 21, doi:10.1145/1480945.1480950. URL <http://portal.acm.org/citation.cfm?doid=1480945.1480950>.
 29. Oracle. Throwable (Java Platform SE6) 2011. URL <http://docs.oracle.com/javase/6/docs/api/java/lang/Throwable.html>.
 30. De Lucia A. Program slicing: Methods and applications. *International Working Conference on Source Code Analysis and Manipulation*, IEEE Computer Society, 2001; 144.
 31. Holzmann GJ. The model checker SPIN. *IEEE Transactions on software engineering* 1997; **23**(5):279–295.

32. Kim S, Zimmermann T, Nagappan N. Crash Graphs: An Aggregated View of Multiple Crashes to Improve Crash Triage. *International Conference on Dependable Systems and Networks (DSN)*, 2013; 486–493.
33. Opyrchal L, Prakash A. Efficient Object Serialization in Java Lukasz Opyrchal and Atul Prakash. *Proceedings. 19th IEEE International Conference on Distributed Computing Systems*, 1999; 96–101.
34. Xstream. Xstream 2011. URL <http://xstream.codehaus.org/>.
35. Apache Software Foundation. Apache Ant. URL <http://ant.apache.org/>.
36. CollabNet. Tigris.org: Open Source Software Engineering. URL <http://www.tigris.org/>.
37. Wellington B. dnsjava 2013. URL <http://www.dnsjava.org/>.
38. Object Refinery Limited. JFreeChart 2005. URL <http://jfree.org/jfreechart/>.
39. The Apache Software Foundation. Log4j 2 Guide - Apache Log4j 2 1999. URL <http://logging.apache.org/log4j/2.x/>.
40. NASA. Open Mission Control Technologies 2009. URL <https://sites.google.com/site/openmct/>.
41. Apache Software Foundation. Apache PDFBox — A Java PDF Library 2014. URL <http://pdfbox.apache.org/>.
42. Apache Software Foundation. Apache Struts Project 2000. URL <http://struts.apache.org/>.